

PREFACE

Our purpose in writing this second edition—more than a quarter century after the original—remains the same: to give a systematic account of the major topics in pattern recognition, based whenever possible on fundamental principles. We believe that this provides the required foundation for solving problems in more specialized application areas such as speech recognition, optical character recognition, or signal classification. Readers of the first edition often asked why we combined in one book a Part I on pattern classification with a Part II on scene analysis. At the time, we could reply that classification theory was the most important domain-independent theory of pattern recognition, and that scene analysis was the only important application domain. Moreover, in 1973 it was still possible to provide an exposition of the major topics in pattern classification and scene analysis without being superficial. In the intervening years, the explosion of activity in both the theory and practice of pattern recognition has made this view untenable. Knowing that we had to make a choice, we decided to focus our attention on classification theory, leaving the treatment of applications to the books that specialize on particular application domains. Since 1973, there has been an immense wealth of effort, and in many cases progress, on the topics we addressed in the first edition. The pace of progress in algorithms for learning and pattern recognition has been exceeded only by the improvements in computer hardware. Some of the outstanding problems acknowledged in the first edition have been solved, whereas others remain as frustrating as ever. Taken with the manifest usefulness of pattern recognition, this makes the field extremely vigorous and exciting.

While we wrote then that pattern recognition might appear to be a rather specialized topic, it is now abundantly clear that pattern recognition is an immensely broad subject, with applications in fields as diverse as handwriting and gesture recognition, lipreading, geological analysis, document searching, and the recognition of bubble chamber tracks of subatomic particles; it is central to a host of human-machine interface problems, such as pen-based computing. The size of the current volume is a testament to the body of established theory. Whereas we expect that most of our readers will be interested in developing pattern recognition systems, perhaps a few will be active in understanding existing pattern recognition systems, most notably human and animal nervous systems. To address the biological roots of pattern recognition would of course be beyond the scope of this book. Nevertheless, because neurobiologists and psychologists interested in pattern recognition in the natural world continue to rely on more advanced mathematics and theory, they too may profit from the material presented here.

Despite the existence of a number of excellent books that focus on a small set of specific *techniques*, we feel that there is still a strong need for a book such as ours, which takes a somewhat different approach. Rather than focus on a specific technique such as neural networks, we address a specific *class of problems*—pattern recognition problems—and consider the wealth of different techniques that can be applied to it. Students and practitioners typically have a particular problem and need to know which technique is best suited for their needs and goals. In contrast, books that focus on neural networks may not explain decision trees, or nearest-neighbor methods, or many other classifiers to the depth required by the pattern recognition practitioner who must decide among the various alternatives. To avoid this problem, we often discuss the relative strengths and weaknesses of various classification techniques.

These developments demanded a unified presentation in an updated edition of Part I of the original book. We have tried not only to expand but also to improve the text in a number of ways:

New Material. The text has been brought up to date with chapters on pattern recognition topics that have, over the past decade or so, proven to be of value: neural networks, stochastic methods, and some topics in the theory of learning, to name a few. While the book continues to stress methods that are statistical at root, for completeness we have included material on syntactic methods as well. “Classical” material has been included, such as Hidden Markov models, model selection, combining classifiers, and so forth.

Examples. Throughout the text we have included worked examples, usually containing data and methods simple enough that no tedious calculations are required, yet complex enough to illustrate important points. These are meant to impart intuition, clarify the ideas in the text, and to help students solve the homework problems.

Algorithms. Some pattern recognition or learning techniques are best explained with the help of algorithms, and thus we have included several throughout the book. These are meant for clarification, of course; they provide only the skeleton of structure needed for a full computer program. We assume that every reader is familiar with such pseudocode, or can understand it from context here.

Starred Sections. The starred sections (*) are a bit more specialized, and they are typically expansions upon other material. Starred sections are generally not needed to understand subsequent unstarred sections, and thus they can be skipped on first reading.

Computer Exercises. These are not specific to any language or system, and thus can be done in the language or style the student finds most comfortable.

Problems. New homework problems have been added, organized by the earliest section where the material is covered. In addition, in response to popular demand, a Solutions Manual has been prepared to help instructors who adopt this book for courses.

Chapter Summaries. Chapter summaries are included to highlight the most important concepts covered in the rest of the text.

Graphics. We have gone to great lengths to produce a large number of high-quality figures and graphics to illustrate our points. Some of these required extensive

calculations, selection, and reselection of parameters to best illustrate the concepts at hand. Study the figures carefully! The book's illustrations are available in Adobe Acrobat format that can be used by faculty adopting this book for courses to create presentations for lectures. The files can be accessed through a standard web browser or an ftp client program at the Wiley STM ftp area at:

ftp://ftp.wiley.com/public/sci_tech_med/pattern/

The files can also be accessed from a link on the Wiley Electrical Engineering software supplements page at:

http://www.wiley.com/products/subject/engineering/electrical/software_supplem_elec_eng.html

Mathematical Appendixes. It comes as no surprise that students do not have the same mathematical background, and for this reason we have included mathematical appendixes on the foundations needed for the book. We have striven to use clear notation throughout—rich enough to cover the key properties, yet simple enough for easy readability. The list of symbols in the Appendix should help those readers who dip into an isolated section that uses notation developed much earlier.

This book surely contains enough material to fill a two-semester upper-division or graduate course; alternatively, with careful selection of topics, faculty can fashion a one-semester course. A one-semester course could be based on Chapters 1–6, 9 and 10 (most of the material from the first edition, augmented by neural networks and machine learning), with or without the material from the starred sections.

Because of the explosion in research developments, our historical remarks at the end of most chapters are necessarily cursory and somewhat idiosyncratic. Our goal has been to stress important references that help the reader rather than to document the complete historical record and acknowledge, praise, and cite the established researcher. The Bibliography sections contain some valuable references that are not explicitly cited in the body of the text. Readers should also scan through the titles in the Bibliography sections for references of interest.

This book could never have been written without the support and assistance of several institutions. First and foremost is of course Ricoh Innovations (DGS and PEH). Its support of such a long-range and broadly educational project as this book—amidst the rough and tumble world of industry and its never-ending need for products and innovation—is proof positive of a wonderful environment and a rare and enlightened leadership. The enthusiastic support of Morio Onoe, who was Director of Research, Ricoh Company Ltd. when we began our writing efforts, is gratefully acknowledged. Likewise, San Jose State University (ROD), Stanford University (Departments of Electrical Engineering, Statistics and Psychology), The University of California, Berkeley Extension, The International Institute of Advanced Scientific Studies, the Niels Bohr Institute, and the Santa Fe Institute (DGS) all provided a temporary home during the writing of this book. Our sincere gratitude goes to all.

Deep thanks go to Stanford graduate students Regis Van Steenkiste, Chuck Lam and Chris Overton who helped immensely on figure preparation and to Sudeshna Adak, who helped in solving homework problems. Colleagues at Ricoh aided in numerous ways; Kathrin Berkner, Michael Gormish, Maya Gupta, Jonathan Hull

and Greg Wolff deserve special thanks, as does research librarian Rowan Fairgrove, who efficiently found obscure references, including the first names of a few authors. The book has been used in manuscript form in several courses at Stanford University and San Jose State University, and the feedback from students has been invaluable. Numerous faculty and scientific colleagues have sent us many suggestions and caught many errors. The following such commentators warrant special mention: Leo Breiman, David Cooper, Lawrence Fogel, Gary Ford, Isabelle Guyon, Robert Jacobs, Dennis Kibler, Scott Kirkpatrick, Daphne Koller, Benny Lautrup, Nick Littlestone, Amir Najmi, Art Owen, Rosalind Picard, J. Ross Quinlan, Cullen Schaffer, and David Wolpert. Specialist reviewers—Alex Pentland (1), Giovanni Parmigiani (2), Peter Cheeseman (3), Godfried Toussaint (4), Padhraic Smyth (5), Yann Le Cun (6), Emile Aarts (7), Horst Bunke (8), Tom Dietterich (9), Anil Jain (10), and Rao Vemuri (Appendix)—focused on single chapters (as indicated by the numbers in parentheses); their perceptive comments were often enlightening and improved the text in numerous ways. (Nevertheless, we are responsible for any errors that remain.) George Telecki, our editor, gave the needed encouragement and support, and he refrained from complaining as one manuscript deadline after another passed. He, and indeed all the folk at Wiley, were extremely helpful and professional. Finally, deep thanks go to Nancy, Alex, and Olivia Stork for understanding and patience.

DAVID G. STORK
 RICHARD O. DUDA
 PETER E. HART

Menlo Park, California
August, 2000

PATTERN CLASSIFICATION
